

Exam Code: 000-716

Exam Name: DB2 Data Warehouse Edition V9.1

Vendor: IBM

Version: DEMO

Part: A

1: A classification model (in PMML format) was created by a third party vendor tool such as SPSS. After importing the model into DB2, where is it stored?

A. In a user-specified table <schema>.<tablename> under an automatically generated name, e.g. SPSS_017.

B. In the table IDMMX.CLASSIFMODELS under a user-specified MODELNAME. The default is the value of the attribute "modelname" in the PMML TreeModel.

C. In the table <userid>.CLASSIFMODELS under a user-specified MODELNAME. The default is the value of the attribute "modelname" in the PMML TreeModel.

D. In a user-specified table <schema>.<tablename> under a user-specified MODELNAME. The default is the value of the attribute "modelname" in the PMML TreeModel.

Correct Answers: B

2: Which of the following must be fulfilled to execute a mining flow that creates a classification model in the database?

A. All input and output ports of all operators must be connected.

B. The flow must contain a Predictor operator whose properties define a Cost Matrix.

C. The flow must contain a Predictor operator whose output port must be connected to a Visualizer.

D. The flow must contain a Predictor operator whose properties (Mining Settings) define a categorical target column.

Correct Answers: D

3: Application data changes over time and differs from that used to train a classification model M. Assume that a mining flow with a Tester operator with M and the application data as input and a test model T as output has been defined. What must be done to check if M can still safely be applied?

A. Open T in a visualizer and check the tree depth. If the tree depth is greater than 10, the model is OK.

B. Open T in a visualizer and check the model quality figures. If the quality is greater than 0.5, the model is OK.

C. Open M and T in a visualizer and check the tree depth. If the tree depth of T is greater than the tree depth of M, the model is OK.

D. Open M and T in a visualizer and check the model quality figures. If the quality is about the same as the quality of M, the model is OK.

Correct Answers: D

4: As part of a market basket analysis effort, a retail company needs to analyze sales transaction data to understand which products tend to be bought together by the same customer during each visit to the store. Which data mining function is most appropriate for this business problem?

A. Clustering

B. Regression

C. Associations

D. Classification

Correct Answers: C

5: A group of Alphablox users needs a report which contains rule heads and bodies of an Association model in a tabular format. Which of the following describes a complete Mining Flow that loads the model output of an Associations operator into a table?

- A. Table Source --> Associations --> Table Target
- B. Table Source --> Associations --> Associations Extractor
- C. Table Source --> Associations --> Visualizer --> Table Target
- D. Table Source --> Associations --> Associations Extractor --> Table Target

Correct Answers: D

6: An investment banking company wants to use data mining to predict customer lapsing behavior. The cost of losing existing customers is very high, and the company wants to keep all customers happy. Based on a recent customer survey with a small subset of existing customers, 2% of customers indicated they were not satisfied. To predict whether or not a customer is satisfied, Predictor and Tester operators in DWE Design Studio were used to analyze the survey results. Which of the following should be done in this situation?

- A. Increase the Maximum tolerated systematic error (RSquared).
- B. Decrease the Maximum tolerated systematic error (RSquared).
- C. Increase the weight of "Satisfied (Actual) & Dissatisfied (Predicted)" in the Cost Matrix parameter option.
- D. Increase the weight of "Dissatisfied (Actual) & Satisfied (Predicted)" in the Cost Matrix parameter option.

Correct Answers: D

7: An insurance company wants to use data mining to predict fraudulent claims. Based on historical data, 1% of claim transactions were classified as fraud. To work on this problem, the Predictor and Tester mining operators in DWE Design Studio were used to analyze a small subset of claim history data. Which preprocessing function operator is the most efficient mechanism to feed data into the Predictor mining operator?

- A. Sampler
- B. Select List
- C. Random Splitter
- D. Field Extractor

Correct Answers: C

8: Which mining model needs to be processed through training and testing phase before being applied into the Scoring function?

- A. Sequential
- B. Clustering
- C. Association
- D. Classification

Correct Answers: D

9: Given the following association rule:

[Tomato][Carrot] \Rightarrow [Orange Juice] Support = 10 Confidence = 80 Lift = 4

Which of the following can be deduced from this information?

- A.The number of orange juice sales is made up of 20% of all sales transactions.
- B.Customers buy the combination of tomato, carrot and orange juice 80% more than expected.
- C.There is a 10% chance that customers also buy orange juice when they buy tomato with carrot.
- D.Out of 2 million sales transactions, approximately 1,600,000 show that tomato, carrot and orange juice are bought together.

Correct Answers: A

10: Which of the following mining operator provides a model output that can be directed into the Visualizer operator in the DWE Mining Flow.

- A.Associations
- B.Cluster Table
- C.Predict Column
- D.Find Deviations

Correct Answers: A

11: SQL code is generated from a (correct) mining flow containing Preprocessing operators, a Clustering operator and a Visualizer. Then the code is copied to a file MINING.SQL in order to execute it in a DB2 command window using "db2 -tf MINING.SQL". Which of the following statements is true?

- A.The statement "DB2 CONNECT TO <database>" must be issued first.
- B.Everything runs to completion and the visualizer is launched in the end.
- C.Preprocessing works, but the Clustering operator cannot be executed from command line.
- D.The Preprocessing operators will fail, because the user will not have the authorization to create tables and/or views.

Correct Answers: A

12: After having tested a particular project using the Design Studio, the administrator is to now deploy the application into a production environment. The administrator will use a Design Studio wizard to do which of the following?

- A.Create a local schedule and deploy the application locally.
- B.Connect to the production system and deploy the application directly.
- C.Export the project in an XML format in order to be imported into the production system.
- D.Create an enterprise archive file that will later be installed in WebSphere Application Server.

Correct Answers: D

13: The use of the Pass through property for a data station allows:

- A.the data station to be ignored.
- B.debugging of the data passed into the data station.
- C.data to be passed into the data station as a temporary table.
- D.data to be passed from an input operation in order to be externalized.

Correct Answers: A

14: A data warehouse design administrator is using the Design Studio to create a project to move data from a designated source to a target. The administrator finds the need to have some steps execute in an iterative fashion in which the number of iterations may vary. How should this be defined?

A.Create a data flow that employs a subflow and encapsulates that subflow within an iterative operation.

B.Create a subflow of the required steps and place them in a data flow based upon the desired number of iterations.

C.Create a control flow that employs a subflow that contains the desired steps and encapsulate the subflow within an iterative operation.

D.Create a control flow that employs a data flow that contains the desired steps and encapsulate the data flow within an iterative operation.

Correct Answers: D

15: An administrator is building a data flow and has the need to join three tables together. The Table Join operator is dragged on to the work area and only has two input ports defined. Which of the following is the recommended technique to follow?

A.Use the SQL Merge operator to join all three tables.

B.Add an additional input port to the Table Join operator.

C.Use the Table Join operator to join two of the tables and then use the properties view to change the generated SQL.

D.Use the Table Join operator to join two of the tables and then use a second Table Join operator to outer join the third.

Correct Answers: B

16: In order to import a data model from ERWIN into DWE 9.1, which steps are needed?

A.Connect to the ERWIN data model from the DWE Administration Console and import the model.

B.Export the model from ERWIN as an XML file and import the XML file into the DWE Design Studio.

C.Export the data model from ERWIN to a DDL file, and use the DWE Design Studio to reverse engineer the model from the DDL file.

D.Export the data model from ERWIN to a DDL file, import the DDL file into the database, and use the DWE Administration Console to mark the imported tables to be used in the model.

Correct Answers: C

17: An administrator is using the Design Studio to create a Control Flow and needs to process a number of the same steps for each file. The number and names of the files that will be processed may vary on a monthly basis. For this to occur in an efficient manner, which of the following would the administrator do?

A.Pass to the iterator the name of a directory that contains the files to be processed.

B.Put the names of the files as rows in a DB2 table and pass the DB2 table to the iterator.

C.Define a variable that contains the names of all of the files to be processed and pass that

variable to the iterator.

D.Setup an iterator that will accept a variable that will determine the number of iterations and have a variable for each file passed into the iteration.

Correct Answers: A

18: Which of the following SQL Warehousing Tool flows can be reused within a project?

A.Subflow

B.Data Flow

C.Mining Flow

D.Control Flow

Correct Answers: A

19: An administrator is using the Design Studio to move data from a source to a target. Which of the following is true?

A.Data flows are used to schedule subflows.

B.Data flows are reusable units of work within a control flow's subflow.

C.Data flows are combined together and scheduled to run via a control flow.

D.Data flows are the smallest unit of work which can be independently scheduled in a control flow.

Correct Answers: C

20: A data warehouse project application that is to perform weekly updates has been prepared for deployment. What steps are needed to complete the deployment and ensure the weekly updates?

A.Use the DWE Design Studio to specify data sources, deploy the application, and schedule the weekly updates.

B.Use the DWE SQL Warehousing portion of the DWE Administration Console to specify data sources, deploy the application, and schedule the weekly updates.

C.Use the DWE Design Studio to deploy the application, then use the DWE Administration Console to identify the needed data sources and to schedule the application.

D.Use the DWE Design Studio to deploy the application and to identify the needed data sources, then use the DWE Administration Console to schedule the application.

Correct Answers: B