



Vendor: Oracle

Exam Code: 1Z0-027

Exam Name: Oracle Exadata X3 and X4 Administration

Version: DEMO

QUESTION 1

Which two statements are true about the use of direct path loads when selecting from external tables in a database on a Database Machine?

- A. INSERT INTO . . . SELECT FROM statements, executed serially, which select from external tables, require the APPEND hint to use direct path loading.
- B. CREATE TABLE . . . AS SELECT statements, which select from external tables, attempt to use in direct path loading automatically.
- C. CREATE TABLE . . . AS SELECT statements, which select from external tables, require the APPEND hint to use direct path loading.
- D. INSERT INTO . . . SELECT FROM statements, executed serially, which select from external tables, are unable to use direct path loading.

Answer: AB

Explanation:

A CTAS (Create table as select) will always use direct path (B, not C) load but IAS (Insert as select) statement will not. In order to achieve direct path load with an IAS statement you must add the APPEND hint to the command (A, not D).

Direct path loads can also run in parallel. You can set the parallel degree for a direct path load either by adding the PARALLEL hint to the CTAS or IAS statement or by setting the PARALLEL clause on both the external table and the table into which the data will be loaded.

Once the parallel degree has been set at CTAS will automatically do direct path load in parallel but an IAS will not. In order to enable an IAS to do direct path load in parallel you must alter the session to enable parallel DML.

Note:

*** Parallel Direct Path Load**

The key to good load performance is to use direct path loads wherever possible. A direct path load parses the input data according to the description given in the external table definition, converts the data for each input field to its corresponding Oracle data type, then builds a column array structure for the data. These column array structures are used to format Oracle data blocks and build index keys. The newly formatted database blocks are then written directly to the database, bypassing the standard SQL processing engine and the database buffer cache.

Reference: Best Practices for Implementing a Data Warehouse on the Oracle Exadata Database Machine; Using CTAS & Exchange Partition Replace IAS for Copying Partition on Exadata

QUESTION 2

Your customer wants you to partition the database and storage grids in his X3-2 full rack, creating database clusters and two storage grids.

One cluster will be used for production and should consist of 6 database servers and 11 cells from the first storage grid.

The other cluster will be used for test and development, and should consist of 2 database servers and 3 cells from the second storage grid.

The storage must be partitioned so that the cells are visible only to the appropriate database servers based on the description above.

What must be done to achieve this?

- A. Configure Exadata realms using Oracle ASM scoped security mode.
- B. Configure Exadata realms using Database scoped security mode.
- C. Edit the CELLIP.ORA file on each database server to contain IP addresses of cells in the storage grid associated with cluster to which that database server belongs.
- D. Edit the CELLINIT.ORA file on each database server to contain IP addresses of cells in the storage grid associated with the cluster to which that database server belongs.
- E. Edit the CELLIP.ORA file on each database server to contain IP addresses of database servers

which are allowed access to specific cells in the same storage grid.

- F. Edit the CELLIP.ORA file on each cell to contain IP addresses of database servers in the database server grid that are associated with the storage grid to which that cell belongs.

Answer: C

Explanation:

cellip.ora

The cellip.ora is the configuration file, on every compute node, that tells ASM instances which cells are available to this cluster.

Here is a content of a typical cellip.ora file for a quarter rack system:

```
$ cat /etc/oracle/cell/network-config/cellip.ora
```

```
cell="192.168.10.3"
```

```
cell="192.168.10.4"
```

```
cell="192.168.10.5"
```

Now that we see what is in the cellip.ora, the grid disk path, in the examples above, should make more sense.

Note:

- * cellinit.ora decides which network takes storage traffic.

- * cellip.ora - list of cells, new cells can be added dynamically without shutdown

QUESTION 3

Which three statements are true about the initial storage configuration after the standard deployment of a new database Machine?

- A. The Data_<DBM_Name> and RECO_<DBM_Name> ASM diskgroups are built on part of the interleaved griddisks.
- B. The Data_<DBM_Name> and RECO_<DBM_Name> ASM diskgroups are built on top of the non- interleaved griddisks.
- C. There is a free space available on the hard disks inside the database servers for possible extension of the /u01 file system.
- D. There is no free disk space available on the hard disk inside the database servers for possible extensions of the /u01 file system.
- E. There is free space available on the flashdisk inside the Exadata storage servers for possible use as flash-based griddisks.
- F. There is no free disk available on the flashdisk inside the Exadata storage servers for possible use as flash-based griddisks.

Answer: BCF

Explanation:

B(not A): Griddisks are the fourth layer of abstraction, and they will be the Candidate Disks to build your ASM diskgroups from. By default (interleaving=none on the Celldisk layer), the first Griddisk that is created upon a Celldisk is placed on the outer sectors of the underlying Harddisk. It will have the best performance therefore. If we follow the recommendations, we will create 3 Diskgroups upon our Griddisks: DATA, RECO and SYSTEMDG.

Note:

- * non-root partition (/u01).

- * Oracle Database files: DATA disk group.

- * Flashback log files, archived redo files, and backup files: RECO disk group.

QUESTION 4

You are in the planning stage of the network configuration for your Database Machine. The requirements are:

1. A fault-tolerant network, providing higher availability for connections to database instances
2. Fault tolerance providing higher availability for connections to perform management functions on the database and storage servers.
3. Full monitoring of all Database Machine components using Enterprise Manager

Which three components require external Ethernet network cables to connect your existing network infrastructure to your database machine to satisfy this requirement?

- A. Database servers
- B. Exadata storage servers
- C. InfiniBand storage servers.
- D. Power distribution units (PDUs)
- E. Oracle Ethernet switch

Answer: ABC

QUESTION 5

Identify two permitted uses of external InfiniBand connections to a Database Machine.

- A. To connect an ExaLogic Elastic Cloud Machine
- B. To monitor the InfiniBand network using Enterprise Manager
- C. To connect an external tape library
- D. To use a bonded client access network
- E. To use a bonded management network

Answer: AB

Explanation:

A: Combining the Oracle Exalogic Elastic Cloud with the Oracle Exadata Database Machine for SAP NetWeaver

The InfiniBand fabric that spans Exalogic and Exadata components provides the following key ways of simplifying and accelerating SAP NetWeaver installations running on Exalogic Note: A high specialized database networking protocols connects all the components inside an Exadata Database Machine. External connectivity to the Exadata Database Machine is provided through standard 1 Gigabit and 10 Gigabit Multiple X3 even larger configurations.

QUESTION 6

You are evaluating the performance of a SQL statement that accesses a very large table, and have run the following query producing the output shown:

```
SQL> SELECT s.name, m.value/1024/1024 MB FROM V$SYSSTAT s, V$MYSTAT m
2 WHERE s.statistic# = m.statistic# AND
3 (s.name LIKE 'physical%total bytes' OR s.name LIKE 'cell phys%'
4 OR s.name LIKE 'cell IO%');
```

NAME	MB
physical read total bytes	19047.2266
physical write total bytes	0
cell physical IO interconnect bytes	4808.85828
cell physical IO bytes pushed back due to excessive CPU on cell	0
cell physical IO bytes saved during optimized file creation	0
cell physical IO bytes saved during optimized RMAN file restore	0
cell physical IO bytes eligible for predicate offload	18005.6953
cell physical IO bytes saved by storage index	0
cell physical IO interconnect bytes returned by smart scan	3767.32703
cell IO uncompressed bytes	18005.6953

For which two reasons would the; "physical read total bytes" statistic be greater than the "cell physical IO bytes eligible for predicate offload" statistic?

- A. There is an index on the column used in the where clause, causing "cell multiblock physical reads" to be requested by the database instance, resulting in additional I/O for blocks in the cells.
- B. The table is an IOT and has an overflow segment, causing "cell multiblock physical reads" to be requested by the database instance, resulting in additional I/O for block in the cells.
- C. There is an uncommitted transaction that has modified some of the table blocks, causing some "cell single block physical reads" to be requested by the database instance, resulting in additional I/O for block in the cells.
- D. The table is an indexed clustered table, causing "cell single block physical reads" to be requested by the database instance, resulting in additional I/O for blocks in the cells.
- E. There are migrated rows in the table, causing some "cell single block physical reads" to be requested by the database instance, resulting in additional I/O for blocks in the cells.

Answer: BE

Explanation:

Note:

* physical read total bytes: the size of the segment to read is known by the database, and must be read entirely from the database's perspective.

* cell physical IO bytes eligible for predicate offload: this statistic shows the amount of data which the cell server is able to process on behalf of the database, instead of the database processing and the cell server just delivering blocks.

* Cell physical IO bytes eligible for predicate offload --- This number should be high The higher the number more MB/GB is filtered out at the cell level itself rather sending it to the buffer cache to filter the rows.

* In this case, all bytes are processed on the cellserver (cell physical IO bytes eligible for predicate offload=physical read total bytes)

* Cell Offloading:-

The storage cells are intelligent enough to process some workload inside them, saving the database nodes from that work. This process is referred to as cell offloading.

QUESTION 7

You must apply patches and patch bundles in a rolling fashion, if possible, on the components of your Database Machine.

You use RAC for your database and also use Data Guard, having standby database on another Database Machine.

You wish to have scripts that contain the appropriate command to patch your environment. Your patch bundle is on the first database server and is located at /u01/stage. You have downloaded the opplan utility to the first database server and run the following:

\$ORACLE_HOME/oplan/oplangenerateApplySteps /u01/stage Which two are true concerning opplan?

- A. It will generate instructions for patching only storage servers.
- B. It will generate instructions for patching all the components on the Database Machine.
- C. It will generate instructions for patching the RAC Oracle Home on the primary location.
- D. It will generate instructions for patching the Data Guard environment on the standby location.
- E. It will generate instructions for patching the Grid Infrastructure home on the primary location.

Answer: CE

Explanation:

Oplan generates instructions for all of the nodes in the cluster.

Note:

* Oplan is a utility that facilitates you with the application of bundle patches on Exadata compute nodes via Opatch. This new utility helps you with the patch process by generating step-by-step instructions telling you how to apply a bundle patch in your environment.

* Exadata is much more than a "database-in-a-box." It is a set of compute nodes (think RAC node servers) combined with ultra-fast infiniband (...and 10GB ethernet, and multiple 1GB ethernet) and storage nodes.

Incorrect:

Not D: There is no support for Oracle DataGuard

QUESTION 8

Consider the following list of software components:

1. DCLI
2. Management Server (MS)
3. ASM Instance
4. RDBMS instance
5. Restart Server (RS)
6. Cellcli
7. Cell Server (CELLSRV)
8. Diskmon

Identify the location where these software components may run in the standard Database machine deployment.

- A. 3, 4 and 8 run on the database servers; 1, 2, 5, 6 and 7 run on the Exadata Storage servers.
- B. 4 and 8 run on the database servers; 1, 2, 3, 5, 6 and 7 run on the Exadata Storage servers.
- C. 1, 3 and 4 run on the database servers; 2, 5, 6, 7 and 8 run on the Exadata Storage servers.
- D. 3, 4 and 8 run on the database servers; 1, 2, 5, 6 and 7 run on the Exadata Storage servers.
- E. 3, 4 and 8 run on the database servers; 1, 2, 5, 7 and 8 run on the Exadata Storage servers.
- F. 1, 3, 4 and 8 run on the database servers; 1, 2, 5, 7 and 8 run on the Exadata Storage servers.

Answer: B

Explanation:

* (not C or not F): Exalogic includes a tool called DCLI (Distributed Command Line Interface) that can be used to run the same commands on all or a subset of compute nodes in parallel.

* Cellcli (not E): on Exadata Storage Servers.

* RDBMS instance on database server.

* CELLSRV on Exadata Storage Servers.

* ASM on Exadata Storage Servers (not D)

After an Oracle ASM instance has been installed on a single-instance Oracle Database or in an Oracle Real Application Clusters (Oracle RAC) environment, the final Oracle ASM configuration can be performed.

QUESTION 9

You plan to monitor storage servers after configuring an I/O resource manager plan with directives for inter-database plans and intra-database plans. Which two types of metrics would help assess the impact of the intra-database plans on I/O to the storage servers?

- A. Category I/O
- B. Database I/O
- C. Resource Consumer Group I/O

- D. Smart Flash Log I/O
- E. Smart Flash Cache I/O

Answer: BC

Explanation:

B: Database metrics provide information about the size of the I/O load from each database specified in the interdatabase plan.

C: Consumer group metrics provide information about the size of the I/O load from each consumer group specified in a database resource plan. Each database in the interdatabase plan has metrics for each of its consumer groups.

Note:

* I/O Resource Manager (IORM) Settings

Incorrect:

Not A: Category metrics provide information about the size of the I/O load from each category specified in the current IORM category plan.

QUESTION 10

You configuring has two half racks, one with high capacity disks and other with high performance disks and high capacity expansion half rack. There are two RAC clusters, one on each half rack, which have separate storage grids, each consisting of some of the storage servers in the configuration. You are planning your deployment of Enterprise Manager to monitor all the components of this multi- rack Database Machine, and must provide for high availability of the monitoring infrastructure. If the host running the agent which has database machine targets bound to it fails, the monitoring of these targets must be another agent. Which two are true regarding the configuration used to support this?

- A. Enterprise manager support must be deployed to only one Enterprise Manager Agent in each cluster.
- B. A secondary agent must be deployed on a database server in the same cluster as the server hosting the primary agent.
- C. Enterprise Manager support must be deployed to all Enterprise Manager Agents in each cluster.
- D. A secondary agent may be deployed on a database server in a different cluster than the server hosting the primary agent.
- E. A secondary agent must be deployed on a database server in a different cluster than the server hosting the primary agent.
- F. Enterprise Manager support must be deployed to at least two Enterprise Manager Agents in each of the two RAC clusters.

Answer: AD

Thank You for Trying Our Product

Braindump2go Certification Exam Features:

- ★ More than **99,900** Satisfied Customers Worldwide.
- ★ Average **99.9%** Success Rate.
- ★ **Free Update** to match latest and real exam scenarios.
- ★ **Instant Download** Access! No Setup required.
- ★ Questions & Answers are downloadable in **PDF** format and **VCE** test engine format.
- ★ Multi-Platform capabilities - **Windows, Laptop, Mac, Android, iPhone, iPod, iPad**.
- ★ **100%** Guaranteed Success or **100%** Money Back Guarantee.
- ★ **Fast**, helpful support **24x7**.



View list of all certification exams: <http://www.braindump2go.com/all-products.html>



Microsoft



ORACLE



JUNIPER
NETWORKS



EMC²
where information lives[®]

10% Discount Coupon Code: BDNT2014