

Exam Code: 310-345

Exam Name: Sun Certified System Administrator for Sun
Cluster 3.2

Vendor: Sun

Version: DEMO

Part: A

1: You have a ZFS pool containing a single disk c2t0d0. The pool contains a file system that you want to mount on /data. What does the entry look like in the /etc/vfstab file?

- A.ZFS populates the /etc/vfstab entries for you automatically
- B.using ZFS, you do NOT create /etc/vfstab entries
- C./dev/dsk/zpool1/dev/rdisk/zpool1 /data zfs 3 no -
- D./dev/dsk/c2t0d0s0 /dev/rdisk/c2t0d0s0 /data zfs 1 yes global

Correct Answers: B

2: How do you decide what DID name to give to your Veritas Volume Manager disks as you assign them to disk groups?

- A.The cluster assigns the DID name to the VxVM disk.
- B.VxVM automatically assigns the DID name to the disk.
- C.You have to manually assign the DID name to a VxVM disk.
- D.VxVM uses its own naming scheme, it does NOT use the DID name.

Correct Answers: D

3: Which two statements identify the tasks a node can perform, when that node has imported a Veritas disk group? (Choose two.)

- A.The node will increase the size of the private region on the disks within the disk group.
- B.The node physically reads and writes data to the drives within the disk group.
- C.The node waits until all mirrored data is in sync before writting data to the disk group.
- D.The node can voluntarily give up ownership of the disk group by deporting the disk group.

Correct Answers: B D

4: Which two options in /etc/vfstab are valid for the global file system? (Choose two.)

- A.Mount at boot is set to yes.
- B.Mount options must be set to false.
- C.The mount point must have the word global.
- D.The word global is one of the mount options.

Correct Answers: A D

5: How do you define Solaris Volume Manager volumes, when configuring your boot device?

- A.You always define the physical disk partition as a submirror, using the standard c#t#d#s#.
- B.You always define soft partitions, so that you can grow your root file system.
- C.You always define your submirror using the DID device /dev/did/rdsk/d#, so that your volumes are usable in the cluster.
- D.You always make sure both submirrors are attached to the mirror before editing the /etc/vfstab file to mount the mirror devices.

Correct Answers: A

6: Suppose you have a Veritas Volume Manager disk group named appdg with default settings registered using:

```
# cldg create -t vxvm -n node1,node2 appdg.
```

You are now asked to configure the following: appdg should be online on node1 by default, and node1 should take appdg automatically back when it joins the cluster.

Which identifies the correct action to accomplish this task?

- A.No action is required, since this is already the default setup.
- B.cldg set -p preferenced=true -p failback=true appdg
- C.cldg set -p preferenced=node1 -p failback=true appdg
- D.cldg change -p preferenced=true -p failback=true appdg

Correct Answers: B

7: Which is important when installing Sun Cluster with Solaris Volume Manager to manage the boot disks?

- A.The /global/.devices/node@1, /global/.devices/node@2 should have unique volume names for each node in the cluster.
- B.All volume names should be configured to be the same across all cluster nodes.
- C.The /global/.devices/node@1, /global/.devices/node@2 should be configured on a soft partition.
- D.The disk partition 7 on the boot disk should be configured to be 6 Mbytes for the metadbs.

Correct Answers: A

8: You need to take one of your Solaris Volume Manager device groups offline for some emergency repairs. The device group is called ds1 and all devices are used as file systems. You have unmounted all the file systems. Which command would you run to offline the device group?

- A.cldg offline ds1
- B.cldg switch offline
- C.cldg offline -s ds1
- D.cldg switch offline -s ds1

Correct Answers: A

9: Your customer has a fully redundant hardware RAID storage environment. Your customer has also chosen to place the hardware RAID LUNs under Veritas Volume Manager control. Which statement is correct?

- A.Software RAID always has better performance.
- B.You need to add additional RAID protection to the LUNs
- C.Hardware RAID does NOT provide total protection of your data.
- D.You can put the LUNs under Veritas Volume Manager control for the purpose of space management.

Correct Answers: D

10: When would you use a Veritas Volume Manager CDS disk in the bootdg group?

- A.All the time, it is the default.
- B.You cannot use CDS disks in the bootdg group.
- C.When you want all nodes to boot from the same device.
- D.When you want to share your boot device between clusters.

Correct Answers: B

11: You are using Veritas Volume Manager to manage Sun Cluster 3.2 shared storage on a multipathed Hitachi array. Which two software products can you use to manage the multipathing? (choose two)

- A.Veritas DMP
- B.Sun StorEdge Traffic Manager Software (built into Solaris 10)
- C.Hitachi Dynamic Link Manager
- D.Veritas Volume Replicator

Correct Answers: B C

12: When using ZFS for your shared storage within the cluster, the mountpoints default to "/poolname/fsname". You can change the mountpoint name to whatever you want. You want to change the mountpoint for the filesystem myfs1 in the pool devpool to /oradb. Which command do you use?

- A.zset mount=/oradb devpool/myfs1
- B.zpool setmount=/oradb poolname/fsname
- C.zfs /oradb set mountpoint devpool/all
- D.zfs set mountpoint=/oradb devpool/myfs1

Correct Answers: D

13: You just created a Solaris Volume Manager diskset from the shared storage containing two disks, one from each of two arrays. You added two connected nodes as mediators for the diskset. Which two of the following are true? (Choose two.)

- A.You must manually create a partition on the diskset disks to hold a metadb replica.
- B.The metainit commands that you use to build volumes on the disks will fail unless you mirror across arrays.
- C.The mediators allow you to lose a disk and subsequently maintain high availability if a node fails.
- D.If you lose one of the disks and then lose the node that is primary for the diskset, the entire cluster will lose access to the diskset.
- E.The Solaris Volume Manager commands that added disks to the diskset automatically created a single metadb replica on each disk.

Correct Answers: C E

14: You have a four-node cluster, where node1 and node2 are connected to some shared storage, and node3 and node4 are connected to different shared storage.

Now node1 has died and you will not be able to reboot it for some time. Why would you then want to run clq disable node1?

- A.This will prevent node1 from accidentally rebooting until you explicitly run clq enable node1 on one of the remaining nodes.
- B.You may want to run the command but it is actually too late. You can only run the disable option when the node mentioned is still up in the cluster.
- C.This will alter the quorum vote count in the remaining cluster so that if later you also lost node2, the remaining nodes could continue running the cluster.

D.This is simply an informational flag. The output of the clq status will specifically show a maintenance flag for node1, indicating that you recognize it is dead and will not be able to reboot for some time.

Correct Answers: C

15: A customer has a four-node cluster configured as a scalable topology. No quorum devices have been configured. How many nodes must be present in order for the cluster to be able to offer any of its application services?

- A.one
- B.two
- C.three
- D.four

Correct Answers: C

16: A three-node cluster contains three resource groups which have these Nodelists:

one-rg has a Nodelist: node1 node2 node3

two-rg has a Nodelist: node2 node3 node1

three-rg has a Nodelist: node3 node1 node2

The Failback property on each of the resource groups is set to FALSE.

Following a maintenance outage in which node1 and node2 have been shut down and rebooted into the cluster with node1 joining the cluster before node2:

Where will the resource groups be located, and which command should be run to move the resource groups to their preferred nodes?

A.one-rg will be on node1

two-rg will be on node2

three-rg will be on node3

No command will need to be run.

B.one-rg will be on node3

two-rg will be on node3

three-rg will be on node3

clrg remaster +

C.one-rg will be on node1

two-rg will be on node1

three-rg will be on node3

clrg remaster +

D.one-rg will be on node3

two-rg will be on node3

three-rg will be on node3

clrg evacuate -n node3 +

Correct Answers: B

17: You have been asked to change the default 10-second timeout setup of the cluster heartbeat to four seconds. Which identifies the correct procedure to accomplish this task?

A.cluster set -p heartbeat_timeout=4000

- B.cluster change heartbeat_timeout=4000
- C.1) cluster set -p heartbeat_quantum=800
- 2) cluster set -p heartbeat_timeout=4000
- D.1) cluster set -p heartbeat_timeout=400
- 2) cluster set -p heartbeat_quantum=800

Correct Answers: C

18: A cluster administrator wants to give the user pat rights to switch the resource groups in the cluster. Which task should the administrator perform?

- A.Assign the authorization solaris.cluster.admin to the user pat.
- B.Nothing, the authorization is by default part of the Basic Solaris User profile.
- C.roleadd -u 1234 -d /export/home/clusoper -A solaris.cluster.modify usermod -R clusoper pat
- D.Tell pat the URL for the Sun Java Web Console and allow the user to login with their username and password.

Correct Answers: A

19: You have a two-node cluster that is using scsi-2 disk reservations for your shared storage. All your shared disks support scsi-3 and their default fencing property has not been changed. The cluster's global_fencing property is set to pathcount. The quorum device is d7 and you want d7 to remain as your quorum device. Your shared devices are d3-d10. Which set of commands will change all disks to use scsi-3?

- A.cluster set -p global_fencing=scsi3
clq refresh d7
- B.cluster set -p global_fencing=scsi3
clq add d6
clq remove d7
- C.cluster set -p global_fencing=prefer3
clq remove d7
clq add d7
- D.cluster set -p global_fencing=prefer3
clq add d6
clq remove d7
cldev set -p default_fencing=global d7
clq add d7
clq remove d6

Correct Answers: D

20: A patch has been released for a recently discovered security issue. The cluster administrator has downloaded the patch and the README shows that it is a rebooting node patch that must be installed in single-user mode. What steps should be followed to apply this patch with the minimum amount of downtime for the service.

- A.On each node in turn run the commands:
clnode evacuate -n <node>
shutdown -g0 -y -i0

boot -sx

Apply the patch.

Reboot back into cluster.

B.On one node run the following commands:

clrg status

clrg offline +

On each node in turn run the commands:

shutdown -g0 -y -i0

boot -sx

Apply the patch.

Reboot back into cluster.

When all nodes are patched run the following commands from one node:

clrg online +

clrg status

C.Connect to all nodes using the cluster console and apply the patch simultaneously.

Reboot each node in turn.

D.On one node run the commands:

clrg status

cluster shutdown -g0 -y

Reboot each node outside cluster.

boot -sx

Apply the patch on each node and reboot back into cluster.

boot

Correct Answers: A